

# Introduction to online optimization

Yuanzhi Li

Assistant Professor, Carnegie Mellon University

Today

- In the entire course, we have been focusing on how to do optimization in a **stationary environment**.
- **Stationary**: we are give a **fixed** function  $f$ , and want to optimize it.

# This lecture

- We are going to ask the following fundamental question: What if the function  $f$  is actually **changing over time**?
- Changing over time: the function  $f$  is changing as we make a move (for example when we update the current point using gradient descent).
- What can we say in this case?

# This lecture

- This is called online optimization. There are two major categories: **Online learning and Reinforcement learning**. Today we are going to focus on online learning. The next lecture will be about reinforcement learning.
- We will begin by giving the definition, and then we will see some applications.

# Online optimization

- The **online optimization** is the following iterative game:
- At every iteration  $t$ , the **environment (or God)** picks a function  $f_t : \mathcal{D} \rightarrow \mathbb{R}$ .
- The player chooses a point  $x_t \in \mathcal{D}$ , **without knowing  $f_t$** .
- The environment then tells the player **some information about  $f_t$  at  $x_t$** :  $f_t(x_t)$  and/or  $\nabla f_t(x_t)$  and/or  $\nabla^2 f_t(x_t)$  etc.

# Online optimization

- Online optimization is a model of **making decisions in a non-stationary environment**.
- Typical example:
- Playing computer games, the environment changes as the player takes a move. (Reinforcement Learning)
- Pushing news to the users: The news are changing rapidly, probably **much faster than the time required to observe a user's feed back**. (Online Learning)

# Online optimization

- Intuitive difference between online learning and reinforcement learning:
- Typically, in reinforcement learning, the environment is **changing with the player's action**. We want the action to **reinforce** the environmental changes.
- In online learning, typically, the environment is changing **independent of player's action**. (Warning: This is not precise, but it is the easiest way to understand online learning). We want the action to **adapt** to the environment.

# Online optimization

- Main question for online optimization: Can the player minimize the sequence of functions  $\{f_t\}_{t=1}^T$ ?
- Meaning that we want the accumulated loss:

$$\sum_{t \in [T]} f_t(x_t)$$

- to be as small as possible.



# Online optimization

- Key difficult of online optimization:
- The player has to pick a point  $x_t$  before seeing the function  $f_t$ .
- In the pushing news example, we have to push the news to the user and then collect the feedbacks (such as whether the user clicks it or not).

# Online optimization

- On the other hand, stationary optimization can also be viewed as a **special case** of online optimization, where  $f_t = f$ .
- The gradient descent algorithm: Each time the player **picks the point  $x_t$ , and then see  $\nabla f(x_t)$** , and then update

$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

# Online learning

- Today we are going to learn the setting called **online learning**, where we want to minimize

$$\sum_{t \in [T]} f_t(x_t)$$

- So that it can be as small as

$$\min_{x \in \mathcal{D}} \sum_{t \in [T]} f_t(x)$$

- In other words, although  $f_t$  is changing, we only want to do **as good as using the best fixed point  $x$** .
- This is because in online learning, intuitively, the function  $f_t$ 's are “independent of” the player's choice  $x_t$ . (Warning: not precise!)
- In **reinforcement learning**, the goal is different.

- We define the regret of the player as:

$$R := \frac{1}{T} \sum_{t \in [T]} f_t(x_t) - \min_{x \in \mathcal{D}} \frac{1}{T} \sum_{t \in [T]} f_t(x)$$

- We want  $R \rightarrow 0$  (or  $R < 0$ ) as  $T \rightarrow +\infty$ .
- Question: Is it doable?

# Online learning: Applications

- Recall the regret:

$$R := \frac{1}{T} \sum_{t \in [T]} f_t(x_t) - \min_{x \in \mathcal{D}} \frac{1}{T} \sum_{t \in [T]} f_t(x)$$

- Special case 1: When  $f_t = f$  is a convex function, then this is doable using **gradient descent**.
- Special case 2: **Stochastic gradient descent**: When  $f_t$  is a **randomly sampled function** such that for every  $x$ ,

$$\mathbb{E}[f_t(x)] = f(x)$$

- For example, in ERM problem,  $f(W) = \frac{1}{N} \sum_{i \in [N]} \ell(h(W, x_i), y_i)$  and each  $f_t(W) = \ell(h(W, x_j), y_j)$  for  $j$  uniformly sampled from  $[N]$ .
- This is doable using **stochastic gradient descent**.

# Online learning: Application

- Other applications mainly include making decisions online, for example **online portfolio selection (stock market), selling ads, buying lottery tickets etc.**
- Here, at each time we are going to make a decision  $x_t$ , and later we will observe its pay-off. But at that time, the environment ( $f_t$ ) may already changed.

# Online learning: Algorithm

- How to minimize the regret:

$$R := \frac{1}{T} \sum_{t \in [T]} f_t(x_t) - \min_{x \in \mathcal{D}} \frac{1}{T} \sum_{t \in [T]} f_t(x)$$

- When  $f_t = f$  is a convex function, we know how to do it using **gradient descent**.
- Key theory for today: When  $f_t$ 's are convex functions, **gradient descent (or mirror descent in general)** also works!
- Without assuming anything (such as  $f_t$ 's are related or changing slowly etc).

# Online learning: Algorithm

- Gradient descent for online learning:
- At each iteration  $t$ , update:

$$x_{t+1} = x_t - \eta \nabla f_t(x_t)$$

- Theorem: when each  $f_t$  is **convex and  $L$ -Lipschitz**, let  $x^* = \operatorname{argmin}_{x \in [T]} f_t(x)$ . Using learning rate  $\eta = \frac{\|x - x_0\|}{L\sqrt{T}}$ , the regret is bounded by:

$$\frac{1}{T} \sum_{t \in [T]} f_t(x_t) - \frac{1}{T} \sum_{t \in [T]} f_t(x^*) \leq \frac{2\|x^* - x_0\|_2 L}{\sqrt{T}}$$



# Online learning: Algorithm

- The theorem says that for sufficiently large  $T$ ,

$$\frac{1}{T} \sum_{t \in [T]} f_t(x_t) - \min_x \frac{1}{T} \sum_{t \in [T]} f_t(x) \propto \frac{1}{\sqrt{T}}$$

- Why this is true? There might be no connection between each  $f_t$ .
- Case 1: All  $f_t$  are sort of random, then there is no  $x$  such that  $\sum_{t \in [T]} f_t(x)$  is small. **So our baseline is pretty bad.**
- Case 2: Most of the  $f_t$  has a common minimizer  $x^*$ , then since  $f_t$  is convex,  **$-\nabla f_t(x_t)$  is pointing to  $x^*$** : The gradient information is useful!

# Online learning: Algorithm

- We also have the **projected gradient descent** for constraint online learning:
- At each iteration  $t$ , update:

$$x_{t+1} = \Pi_{\mathcal{D}}(x_t - \eta \nabla f_t(x_t))$$

# Online learning: Proof

- Let us directly show the **regret bound** for the projected gradient descent algorithm:
- Recall we defined the gradient mapping:

$$g(x_t) := \frac{1}{\eta}(x_t - \Pi_{\mathcal{D}}(x_t - \eta \nabla f_t(x_t)))$$

- Recall when  $f_t$  is  $L$ -Lipschitzness,  $\|g(x_t)\|_2 \leq \|\nabla f(x_t)\|_2 \leq L$ .
- Recall the **three-term Mirror Descent Lemma** for **projected gradient descent** is given by: for every  $x \in \mathcal{D}$

$$f_t(x_t) \leq f_t(x) + \frac{1}{2\eta} (\|x - x_t\|_2^2 - \|x - x_{t+1}\|_2^2 + 2\eta^2 L^2)$$

# Online learning: Proof

- Now we have: for every  $t$ ,

$$f_t(x_t) \leq f_t(x) + \frac{1}{2\eta} (\|x - x_t\|_2^2 - \|x - x_{t+1}\|_2^2 + 2\eta^2 L^2)$$

- Sum them up from  $t = 1, 2$  up to  $T$ , we have:

$$\sum_{t \in [T]} f_t(x_t) \leq \sum_{t \in [T]} f_t(x) + \sum_{t \in [T]} \frac{1}{2\eta} (\|x - x_t\|_2^2 - \|x - x_{t+1}\|_2^2 + 2\eta^2 L^2)$$

- Which is:

$$\frac{1}{T} \sum_{t \in [T]} f_t(x_t) \leq \frac{1}{T} \sum_{t \in [T]} f_t(x) + \frac{1}{2\eta T} \|x - x_0\|_2^2 + \eta L^2$$

- Picking  $\eta = \frac{\|x - x_0\|}{L\sqrt{T}}$  we complete the proof.

- The **spirit** of Mirror Descent: When  $f_t(x_t)$  is much larger than  $f_t(x)$ , then  $x_t$  is moving closer to  $x$ .
- Thus,  $\nabla f_t(x_t)$  is “useful” to find the optimal point  $x^*$ , without any particular assumption on  $f_t$  (except convexity).

# Online learning application: Non-convex optimization

- The **spirit** of this proof is extremely important, it can be extend to analyzing the **gradient descent** update for **non-convex ERM problem**:
- Suppose we want to minimize  $f(W) = \frac{1}{N} \sum_{i \in [N]} \ell(h(W, x_i), y_i)$  for a convex loss  $\ell$  and a **non-convex model**  $h(W, x)$  (such as a neural network).
- Suppose we update it using gradient descent:

$$W_{t+1} = W_t - \eta \nabla f(W_t)$$

- Each time, we can define a convex function  $f_t$ :

$$f_t(W) = \frac{1}{N} \sum_{i \in [N]} \ell(h(W_t, x_i) + \langle \nabla_W h(W_t, x_i), W - W_t \rangle, y_i)$$

# Online learning application: Non-convex optimization

- Each time, we can define a **convex function**  $f_t$ :

$$f_t(W) = \frac{1}{N} \sum_{i \in [N]} \ell(h(W_t, x_i) + \langle \nabla_W h(W_t, x_i), W - W_t \rangle, y_i)$$

- Key observation:

$$f(W_t) = f_t(W_t), \quad \nabla_W f(W_t) = \nabla_W f_t(W_t)$$

- Therefore, we are doing **convex online learning** using **gradient descent**:

$$W_{t+1} = W_t - \eta \nabla f_t(W_t)$$

# Online learning: Multi-arm bandit

- One famous application is the so called the **multi-arm bandit (MAB)** problem, we consider the so called “full observation” case:
- There are  $m$  machines, each time  $t$ , each machine  $i$  has a reward  $\ell_t(i)$ , and we want to “pull” a machine  $i_t \in [m]$  (before knowing  $\ell_t(i)$ ) and collect the reward. **After that, we see all the rewards at time  $t$ .**
- We want to maximize

$$\frac{1}{T} \sum_{t \in [T]} \ell_t(i_t) - \max_{i \in [m]} \frac{1}{T} \sum_{t \in [T]} \ell_t(i)$$

- Algorithm: Maintain a distribution  $p_t$  over  $[m]$ , each step
- (1). Sample  $i_t \sim p_t$ .
- (2). Update: (for  $\ell_t = (\ell_t(i))_{i \in [m]}$ ):

$$p_{t+1} = \Pi_{\text{distributions}}(p_t + \eta \nabla \ell_t)$$

- Intuition: We can define  $f_t(p) = -\langle \ell_t, p \rangle$ .



# Online learning: Multi-arm bandit

- Actually, the most efficient algorithm here is to use **projected mirror descent** (using KL divergence).
- This is “the best distance” to measure difference between distributions.

# Online learning: Other algorithms

- We saw the **gradient descent** algorithm, actually, essentially all the optimization algorithm we have learnt work in the online learning setting.
- These all works: **Mirror Descent/Adagrad/Proximal Gradient Descent/Newton's method/Ellipsoid**.

# Online newton's method

- We also take a special look at the **Online Newton's Method**, or the so called **online BFGS** optimization algorithm:
- At each iteration, update using a **quasi-newton** step:

$$x_{t+1} = x_t - \eta_t B_t \nabla f_t(x_t)$$

- Where the “inverse Hessian” matrix  $B_t$  is given by: for  $s_{t+1} := \eta_t B_t \nabla f_t(x_t)$  and

$$y_t := \nabla f_t(x_t) - \nabla f_{t-1}(x_{t-1})$$

- Update for  $\rho_t = \langle s_t, y_t \rangle^{-1}$ :

$$B_t = (I - \rho_t s_t y_t^\top) B_{t-1} (I - \rho_t y_t s_t^\top) + \rho_t s_t s_t^\top$$

- The intuition of this update is we want to update  $B_t$  such that  $B_t s_t \approx -y_t$ .

# Online learning, other extensions

- There are several extensions to the online learning setting.
- One of the extension is that at each time  $t$ , we can only get **the function value  $f_t(x_t)$**  as the feedback, instead of  $\nabla f_t(x_t)$ .
- This can be quite natural in some applications, where the gradient is hard to obtain.

# Online learning, other extensions

- Recall: In standard optimization, we can estimate  $\nabla f(x)$  using function value: for every  $v$ ,

$$\langle \nabla f(x), v \rangle = \lim_{t \rightarrow 0} \frac{f(x + tv) - f(x)}{t}$$

- We just need  $d$  many such  $v$  to recover the full gradient.
- However, in online learning, after we query a point  $x$  and observe  $f_t(x)$ , the function has changed to  $f_{t+1}$ .
- We might not be able to query  $d$  points for the same function  $f_t$ .

# Online learning, other extensions

- At each time  $t$ , we can only get **the function value  $f_t(x_t)$**  as the feedback, instead of  $\nabla f_t(x_t)$ .
- Key technique: Using stokes formula: for every  $r \geq 0$ , for  $v$  being a randomly sampled unit vector in  $\mathbb{R}^d$  ( $x \in \mathbb{R}^d$ ),

$$\mathbb{E}_v \left[ \frac{d}{r} f(x + rv) v \right] = \nabla \tilde{f}(x)$$

- Where  $\tilde{f}(x) = \int_{\|v\|_2 \leq 1} f(x + rv) dv$
- Thus,  $\frac{d}{r} f(x + rv) v$  is a **stochastic gradient** for function  $\tilde{f}(x)$ .
- Thus, we can query  $x_t + rv$  and update

$$x_{t+1} = x_t - \frac{d}{r} f_t(x_t + rv) v$$

# Online learning, other extensions

- Another major extension assumes that  $f_t$  is not changing very fast, in the sense that

$$V = \sum_{t=2}^T \|f_t - f_{t-1}\|$$

- is bounded by a not so large value.
- The regret can scale with  $\frac{\sqrt{V}}{T}$  in some cases.

# Online learning, other extensions

- Another major extension want to complete with the best sequence of “slow moving”  $x_t^*$ , in the sense that the baseline now becomes:

$$\min_{\{x_t^*\}_{t \in [T]}} \sum_{t \in [T]} f_t(x_t^*)$$

- Such that  $\sum_{t=2}^T \|x_t^* - x_{t-1}^*\|$  is bounded by a not so large value.
- In this setting, sometimes we can achieve competitive ratio  $\text{poly}(d)$ , in the sense that

$$\left( \sum_{t \in [T]} f_t(x_t) + \sum_{t=2}^T \|x_t - x_{t-1}\| \right) \leq \text{poly}(d) \times \min_{\{x_t^*\}_{t \in [T]}} \left( \sum_{t \in [T]} f_t(x_t^*) + \sum_{t=2}^T \|x_t^* - x_{t-1}^*\| \right)$$