

# Convex Optimization 10-725, Lecture 19: Introduction to non-convex optimization: Simulated Annealing, Evolutionary algorithms

Yuanzhi Li

Assistant Professor, Carnegie Mellon University  
Consulting Researcher, Microsoft Research


Today

# Last lecture

- We learnt the Bayesian optimization.

# Last lecture

- “band of possible functions” + “optimistic exploration”.



# This lecture

- We are going to see some other completely different non-convex optimization algorithms.
- They are in [spirit](#) fundamentally different from gradient based optimization or bayesian optimization.

# Simulated Annealing


- We introduce the algorithm **simulated annealing** for **non-convex** optimization.
- This is a type of **evolutionary optimization algorithm**.
- The **spirit** of **evolutionary algorithms** is very very important in the optimization literature.
- We are going to see how they work.

# Simulated Annealing: Motivation

- For a function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  over a constraint set  $\mathcal{D}$ , we can view it as a **distribution over  $\mathcal{D}$** , the density function is given as: (for some  $\lambda \geq 0$ )

$$p(x) \propto e^{-\lambda f(x)}$$

- Key observation: When  $\lambda$  is close to  $= 0$ ,  $p(x)$  is close to a uniform distribution over set  $\mathcal{D}$ .
- When  $\lambda$  is close to  $= +\infty$ , then  $p(x)$  is close to a **distribution uniformly over the global minimizers of  $f$**



# Simulated Annealing: Motivation

- When  $\lambda$  is close to  $+\infty$ ,  $p(x)$  is close to a uniform distribution over the global minimizers of  $f$ .
- If we can sample a point  $x'$  from distribution


$$p(x) \propto e^{-\lambda f(x)}$$

for sufficiently large  $\lambda$ , then with high probability  $f(x')$  would be close the global minimal value of  $f$ .

- Question: How do we sample from  $p(x) \propto e^{-\lambda f(x)}$ ?

# Simulated Annealing: Motivation


- For a good constraint set  $\mathcal{D}$  (for example, the  $\ell_2$  ball), sampling from a uniform distribution over set  $\mathcal{D}$  is **easy**.
- Sampling from  $p(x) \propto e^{-0 \times f(x)}$  is easy.
- Simulated Annealing: Starting from distribution  $p(x) \propto e^{-0 \times f(x)}$ , we **gradually increase the value of  $\lambda$**  to  $+\infty$  and sample from  $p_\lambda(x) \propto e^{-\lambda f(x)}$  at every iteration.
- Here,  $\lambda$  is called the temperature.





# Simulated Annealing: Motivation

- Simulated Annealing: Starting from distribution  $p(x) \propto e^{-0 \times f(x)}$ , we **gradually increase the value of  $\lambda$**  to  $+\infty$  and sample from  $p_\lambda(x) \propto e^{-\lambda f(x)}$  at every iteration.
- Intuition: When  $\lambda_1$  is close to  $\lambda_2$ , sampling from  $p_{\lambda_2}(x)$  using **uniform samples from  $p_{\lambda_1}(x)$  as a warm start** is easier.




# The Metropolis–Hastings algorithm

- Now we answer the main question: How do we sample from a distribution  $P(x) \propto e^{-\lambda_2 f(x)}$  given samples from distribution  $Q(x) \propto e^{-\lambda_1 f(x)}$  **as a warm start?**
- We introduce the **Metropolis–Hastings algorithm**.

# The Metropolis–Hastings algorithm

- The **Metropolis–Hastings algorithm** to sample from a distribution  $P(x)$  given samples from distribution  $Q(x)$  **as a warm start**:
- Initially, sample a point  $x_0$  from distribution  $Q(x)$ .
- At every iteration  $t$ , sample  $y_{t+1} \sim \mathcal{N}(x_t, \sigma^2 I)$  for some  $\sigma > 0$  (to be tuned based on your specific application).
- Define  $\alpha = \min \left\{ \frac{P(y_{t+1})}{P(x_t)}, 1 \right\}$
- Define

$$x_{t+1} = \begin{cases} y_{t+1} & \text{w.p. } \alpha \\ x_t & \text{w.p. } 1 - \alpha. \end{cases}$$



# The Metropolis–Hastings algorithm

- Main theorem: When  $t \rightarrow \infty$ ,  $x_t \rightarrow$  a sample according to distribution  $P(x)$ .
- In other words, the stationary distribution of the Metropolis–Hastings process is  $P(x)$ .
- Proof: Let  $\pi(x)$  be a distribution, let  $p(x | x')$  be the condition probability (density) of arriving at  $x$  by doing one step of the Metropolis–Hastings process starting from  $x'$ ,
- We just need to verify that when  $\pi(x) = P(x)$ ,


$$\pi(x) = \int_{x'} p(x | x') \pi(x') dx'$$

- Consider two cases:  $x' = x$ ,  $x' \neq x$ .

# The Metropolis–Hastings algorithm


- Let  $n(x | x')$  be the probability density of point  $x$  associated with distribution  $\mathcal{N}(x', \sigma^2 I)$ .
- Define  $\alpha_{x,x'} = \min \left\{ \frac{P(x)}{P(x')}, 1 \right\}$ .
- When  $x = x'$ , we know that

$$\Pr[x' = x | x] = \int_z (1 - \alpha_{z,x}) n(z | x) dz$$



# The Metropolis–Hastings algorithm

- Define  $\alpha_{x,x'} = \min \left\{ \frac{P(x)}{P(x')}, 1 \right\}$ .
- When  $x \neq x'$ , we know that  $p(x | x') = \alpha_{x,x'} n(x | x')$ .



- This implies that when  $x' \neq x$ :

$$\int_{x'} p(x | x') \pi(x') dx' = \int_{x'} \alpha_{x,x'} n(x | x') \pi(x') dx'$$

- Together, we have:

$$\begin{aligned} \int_{x'} p(x | x') \pi(x') dx' &= \int_{x'} \alpha_{x,x'} n(x | x') \pi(x') dx' \\ &+ \pi(x) \int_z (1 - \alpha_{z,x}) n(z | x) dz \end{aligned}$$

# The Metropolis–Hastings algorithm

- Together we have:

$$\begin{aligned}\int_{x'} p(x | x') \pi(x') dx' &= \int_{x'} n(x | x') \alpha_{x, x'} \pi(x') dx' \\ &+ \pi(x) \int_z (1 - \alpha_{z, x}) n(z | x) dz\end{aligned}$$

- Key observation:

$$P(x') \alpha_{x, x'} = \min\{P(x), P(x')\} = P(x) \alpha_{x', x}$$

- Which implies that when  $\pi(x) = P(x)$ ,

$$\begin{aligned}\int_{x'} p(x | x') \pi(x') dx' &= \int_{x'} n(x | x') \alpha_{x', x} \pi(x) dx' \\ &+ \pi(x) \int_z (1 - \alpha_{z, x}) n(z | x) dz\end{aligned}$$

# The Metropolis–Hastings algorithm

- Now we have: when  $\pi(x) = P(x)$

$$\begin{aligned}\int_{x'} p(x | x') \pi(x') dx' &= \int_{x'} n(x | x') \alpha_{x',x} \pi(x) dx' \\ &\quad + \pi(x) \int_z (1 - \alpha_{z,x}) n(z | x) dz\end{aligned}$$

- Key observation:  $n(x | x') = n(x' | x)$ .
- Therefore, when  $\pi(x) = P(x)$ :


$$\begin{aligned}\int_{x'} p(x | x') \pi(x') dx' &= \int_{x'} n(x' | x) \alpha_{x',x} \pi(x) dx' \\ &\quad + \pi(x) \int_z (1 - \alpha_{z,x}) n(z | x) dz \\ &= \pi(x) \int_z n(z | x) dz = \pi(x)\end{aligned}$$

- We show that  $P(x)$  is a stationary distribution.



# The Metropolis–Hastings algorithm

- Key idea: using a warm start  $Q$  that is close to  $P$ , Metropolis–Hastings algorithm converges to the distribution  $P$  faster.
- This is also heuristics (but there is a very non-trivial convergence proof when  $P, Q$  are log-concave distributions).



# The Simulate Annealing algorithm


- **Simulated annealing algorithm** to minimize a function  $f$  in a constraint set  $\mathcal{D}$  in general:
- First, sample a point  $x_0$  uniformly at random from  $\mathcal{D}$ , set  $\lambda_0 > 0$  to be sufficiently small.
- At every iteration  $t$ , let  $\lambda_{t+1} = (1 + \eta)\lambda_t$ .
- Run the **Metropolis–Hastings algorithm** (for multiple steps) to sample  $x_{t+1}$  from distribution

$$p_{\lambda_{t+1}}(x) \propto e^{-\lambda_{t+1}f(x)}$$

starting from  $x_t$  as a “warm start”.

# The Simulate Annealing algorithm


- The **spirit** of **simulated annealing algorithm** to minimize a function  $f$ :
- **Evolve** the distribution  $p$  from uniform at random to **supports only on the minimizers** of  $f$ .



- 
- When we run **Metropolis–Hastings algorithm** for sufficiently many steps at every iteration (steps  $\rightarrow +\infty$ ), simulated annealing **converges to the global minimizer of  $f$** .
- However, the number of steps might be very large, there is **no efficient rate guarantee unless  $f$  is convex**.

# Other evolutionary algorithm

- There is an **extremely simple**, and **powerful** alternative **evolutionary algorithm** used in reinforcement learning, called **evolution strategies** (ES).
- Example: MuJoCo controls



# Other evolutionary algorithm

- **Evolutionary strategies** to **maximize** a function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$
- At every iteration  $t$ , randomly sample  $N$  i.i.d. vectors  $\epsilon_1, \dots, \epsilon_N \sim \mathcal{N}(0, I_{d \times d})$ .
- Compute function value using **a standard deviation**  $\sigma_t$ :  
 $f_i = f(x_t + \sigma_t \epsilon_i)$  for every  $i \in [N]$ .
- Update using learning rate  $\eta$ :

$$x_{t+1} = x_t + \eta \frac{1}{N\sigma_t} \sum_{i=1}^N f_i \epsilon_i$$

- The “lucky ones”  $\epsilon_i$  with higher function values  $f_i$  gets **weighted higher**.

- Actually, the ES algorithm is trying to do **stochastic gradient ascent** on the new objective

$$F_t = f * g_t$$

- Where  $g_t$  is the density function of  $\mathcal{N}(0, \sigma_t^2 I)$ ,  $*$  is the **convolution operation**:

$$[f * g](x) = \int_y f(y)g(y - x)dy$$

- By Stoke's formula,


$$\nabla F_t = f * \nabla g_t$$

# ES algorithm

- Actually, the ES algorithm is trying to do **stochastic gradient ascent** on the new objective

$$F_t = f * g_t$$

- Where  $g_t$  is the density function of  $\mathcal{N}(0, \sigma_t^2 I)$
- ES can even escape **local minima** when  $\sigma_t$  is large.



# Evolutionary algorithms

- There are a lot of other evolutionary algorithms.
- The **spirit**: Initially, starting from a **prior distribution over** the set of parameters.
- At every iteration, **evolve** this distribution a little bit towards **higher quality solutions**

